

# The BFGS Optimization Algorithm

Pratham Lalwani

UC Merced

May 13, 2025

# Outline

- 1 Background
- 2 Quasi-Newton Methods
- 3 Rosenbrock Example
- 4 Results
- 5 Conclusion

# Problem Setup

- Given  $f : \mathbb{R}^n \rightarrow \mathbb{R}$ , say we are interested in minimizing the function, which is,

$$\min_{\mathbf{x} \in \mathbb{R}^n} f(\mathbf{x}).$$

- From calculus,  $\nabla f(\mathbf{x}) = \mathbf{0}$  and solve analytically if it can be done.
- This problem arises everywhere especially nowadays with Machine Learning where  $f(\mathbf{x})$  is usually a cost function we are trying to minimize.

# Gradient Descent and Its Limitations

- With no way to compute a analytic solution one might turn a simple algorithm like Gradient Descent.

# Gradient Descent and Its Limitations

- With no way to compute a analytic solution one might turn a simple algorithm like Gradient Descent.
- The next iterate is given by :  $\mathbf{x}_{k+1} = \mathbf{x}_k - \gamma \nabla f(\mathbf{x}_k)$ , where  $\gamma$  is a fixed constant called step size or learning rate.

# Gradient Descent and Its Limitations

- With no way to compute a analytic solution one might turn a simple algorithm like Gradient Descent.
- The next iterate is given by :  $\mathbf{x}_{k+1} = \mathbf{x}_k - \gamma \nabla f(\mathbf{x}_k)$ , where  $\gamma$  is a fixed constant called step size or learning rate.
- Pros: simple, easy to implement and not computationally expensive (per step)

# Gradient Descent and Its Limitations

- With no way to compute a analytic solution one might turn a simple algorithm like Gradient Descent.
- The next iterate is given by :  $\mathbf{x}_{k+1} = \mathbf{x}_k - \gamma \nabla f(\mathbf{x}_k)$ , where  $\gamma$  is a fixed constant called step size or learning rate.
- Pros: simple, easy to implement and not computationally expensive (per step)
- Cons: slow convergence, sensitivity to step size.

# Gradient Descent and Its Limitations

- With no way to compute a analytic solution one might turn a simple algorithm like Gradient Descent.
- The next iterate is given by :  $\mathbf{x}_{k+1} = \mathbf{x}_k - \gamma \nabla f(\mathbf{x}_k)$ , where  $\gamma$  is a fixed constant called step size or learning rate.
- Pros: simple, easy to implement and not computationally expensive (per step)
- Cons: slow convergence, sensitivity to step size.



# Gradient Descent




Figure: Gradient Descent on  $x^2$

# Gradient Descent




Figure: Gradient Descent on  $x^2$

# But wait a minute

Instead of having a fixed Learning Rate what if we have a adaptive learning rate, which is "greedy". If we have,

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha_k \nabla f(\mathbf{x}).$$

# But wait a minute

Instead of having a fixed Learning Rate what if we have a adaptive learning rate, which is "greedy". If we have,

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha_k \nabla f(\mathbf{x}).$$

We define,

$$\phi(\alpha) := f(\mathbf{x}_k - \alpha \nabla f(\mathbf{x})).$$

# But wait a minute

Instead of having a fixed Learning Rate what if we have a adaptive learning rate, which is "greedy". If we have,

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha_k \nabla f(\mathbf{x}).$$

We define,

$$\phi(\alpha) := f(\mathbf{x}_k - \alpha \nabla f(\mathbf{x})).$$

Now we solve the problem,

$$\min_{\alpha \in [0,1]} \phi(\alpha).$$

# But wait a minute

Instead of having a fixed Learning Rate what if we have a adaptive learning rate, which is "greedy". If we have,

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \alpha_k \nabla f(\mathbf{x}).$$

We define,

$$\phi(\alpha) := f(\mathbf{x}_k - \alpha \nabla f(\mathbf{x})).$$

Now we solve the problem,

$$\min_{\alpha \in [0,1]} \phi(\alpha).$$

The good this about this problem is that it is one-dimensional. But we still have to realize the function we are evaluating underneath might be expensive to evaluate.

# Line Search

Now there are tons of techniques to find the minimum of a 1-D function. We shall use the one which is most familiar to us which is bisection method.

# Line Search

Now there are tons of techniques to find the minimum of a 1-D function. We shall use the one which is most familiar to us which is bisection method. But, when should we terminate?



# Line Search

Now there are tons of techniques to find the minimum of a 1-D function. We shall use the one which is most familiar to us which is bisection method. But, when should we terminate?

There is 2 conditions usually that we need to follow,  
Sufficient Decrease:

$$f(\mathbf{x} - \alpha_k \nabla f(\mathbf{x}_k)) \leq f(x_k) + c_1 \alpha_k \|\nabla f\|^2.$$

and curvature condition,

$$\nabla f(x_k + \alpha_k p_k)^T p_k \geq c_2 \nabla f_k^T p_k$$

These two together make the Wolfe conditions of sufficient decrease.

# A picture is worth thousand words




Figure: Strong Wolfe Condition acceptance regions

# Steepest Descent & Newton

- As we know steepest descent has its own problems such as "zig-zag" behaviour as discussed in class. Thus we would like a better method.
- We also learnt about Newton Update for minimizing scalar functions which is given by

$$\mathbf{x}_{k+1} = \mathbf{x}_k - H_f(\mathbf{x}_k)^{-1} \nabla f(\mathbf{x}_k).$$

Where  $H_f$  is the hessian of  $f$  w.r.t  $\mathbf{x}$

- This is computationally quite expensive as it requires solving a linear system which takes  $\mathcal{O}(n^3)$  time to solve. Where  $n$  is problem dimension.

- We would like to develop a Hessian which doesn't take  $O(n^3)$  time to solve.
- Before we do that some notation. All the algorithms we saw before now take the form

$$\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{p}_k.$$

Where  $\mathbf{p}_k$  is a descent direction. We define  $y_k = \nabla f(\mathbf{x}_{k+1}) - \nabla f(\mathbf{x}_k)$  and  $s_k = \mathbf{x}_{k+1} - \mathbf{x}_k$ . We would like to mimic the Newton search direction, which is  $\mathbf{p}_k^{Newton} = -H_f(\mathbf{x}_k)^{-1} \nabla f(\mathbf{x}_k)$ . We would like a approximate hessian  $B_k \approx H_f(\mathbf{x}_k)$  which doesn't take  $O(n^2)$  to compute.

- These lead us to the BFGS update, introduced by Broyden, Fletcher, Goldfarb and Shanno over the course of 4 papers.

# BFGS Derivation Outline

We start by defining a convex quadratic model as at step  $k$  as:

$$m_k(p) = f(\mathbf{x}_k) + \nabla f(\mathbf{x}_k)^\top + \frac{1}{2} \mathbf{p}^\top B_k \mathbf{p}.$$

The unique minimizer of this quadratic is

$$\mathbf{p}_k = -B_k^{-1} \nabla f(\mathbf{x}_k).$$

Now instead of recomputing  $B_{k+1}$  for next iteration we proceed as follows,

We would like to have  $\nabla m_{k+1}(\mathbf{0}) = \nabla f(\mathbf{x}_k)$  and

$\nabla m_{k+1}(-\alpha_k \mathbf{p}_k) = \nabla f(\mathbf{x}_k)$  to provide a good approximation to the objective function  $f$  around those points.

The first one we get for free,

$$\nabla m_{k+1}(\mathbf{0}) = \nabla f(\mathbf{x}_k).$$

The second one simplifies to,

$$B_{k+1} \mathbf{s}_k = \mathbf{y}_k.$$

This is the secant equation for the second derivative.

# BFGS Derivation Outline

The key benefit of BFGS is that it computes the inverse Hessian directly. Which means we directly find  $H_{k+1}$  such that,

$$B_{k+1}s_k = y_k \implies H_k y_k = s_k.$$

We also would like to make the minimal update on  $H_k$  to get  $H_{k+1}$  and as mentioned earlier it is positive definite,

Which results in the following constrained optimization problem,

$$\min_{H \in \mathbb{R}^{n \times n}} \|W^{1/2}(H - H_k)W^{1/2}\|_F \quad \text{subject to } H = H^T \text{ and } Hy_k = s_k.$$

Which gives the formula:  $H_{k+1} = (I - \rho_k s_k y_k^T) H_k (I - \rho_k y_k s_k^T) + \rho_k s_k s_k^T$   
Which is the BFGS update.

Great thing about it is that it is a small rank-2 update and still keeps the matrix positive definite.

# Algorithm 6.1 (BFGS Algorithm)

---

## Algorithm 1 (BFGS Algorithm)

---

**Require:** Given starting point  $x_0$ , convergence tolerance  $\epsilon > 0$ , inverse Hessian approximation  $H_0$ ;


- 1:  $k \leftarrow 0$ ;
  - 2: **while**  $\|\nabla f_k\| > \epsilon$  &  $k < \text{maxIter}$  **do**
  - 3:   Compute search direction
  - 4:    $p_k = -H_k \nabla f_k$ ;
  - 5:   Set  $x_{k+1} = x_k + \alpha_k p_k$  where  $\alpha_k$  is computed from a line search procedure to satisfy the Wolfe conditions;
  - 6:   Define  $s_k = x_{k+1} - x_k$  and  $y_k = \nabla f_{k+1} - \nabla f_k$ ;
  - 7:   Compute  $H_{k+1}$  ;
  - 8:    $k \leftarrow k + 1$ ;
  - 9: **end while**
-

# Rosenbrock Example

- Test function:  $f(x, y) = (a - x)^2 + b(y - x^2)^2$
- Typical parameters:  $a = 1, b = 100$
- Illustrates curved valley and optimization challenge

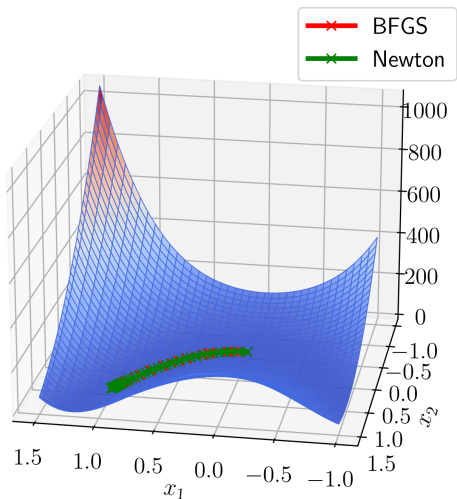


# Rosenbrock Example Results



# Rosenbrock Example Results

Rosenbrock function



# Results

Function	BFGS	Newton	LBFGS	GD
Adjiman Function (2-D)	nan	4.61e-05	5.50e-16	7.70e-01
Rosenbrock N-D (100-D)	9.28e-11	5.13e-04	9.24e-11	2.92e+00
Paviani Function (10-D)	nan	nan	nan	8.72e-02
Csendes Function (10-D)	9.57e-11	1.21e-03	9.16e-11	7.48e-02
Griewank Function (2-D)	nan	1.26e-15	6.65e-11	8.31e-09
Hosaki Function (2-D)	nan	3.17e-05	nan	8.66e-02
Brent Function (2-D)	9.00e-11	2.51e-05	9.90e-11	3.67e+00
Giunta Function (2-D)	2.22e-15	9.40e-05	2.67e-15	1.29e-08
Styblinski-Tang Function (2-D)	2.93e-14	1.12e-04	6.39e-14	9.96e-11
Trid 6 Function (6-D)	nan	2.47e-05	nan	4.08e+00

# Conclusion

- Introduced BFGS as an quasi Newton optimizer.
- Provided description of Wolfe conditions, and an outline of BFGS derivation.
- Compared BFGS against other methods.